

A02 Neural Network Zoo

Course: ITAI2376

Group 4: Oman Malek, Phillip Torres, Tien Manh Nguyen & Lufei Yu

Jan 2026

1. Introduction to Neural Networks

Neural networks are one of the most important concepts in artificial intelligence and deep learning. They are modeled after the human brain and enable computers to perceive patterns, make decisions, and learn from data. A neural network consists of artificial neurons, which receive inputs, perform calculations using mathematical functions, and then produce outputs. Neurons are arranged in layers: the input layer, the hidden layers, and the output layer.

Neural networks improve their performance on tasks such as image and speech recognition, and prediction, by adjusting the connections between neurons during the learning process. All neural networks operate through hierarchical nodes that attempt to identify complex patterns via non-linear activation functions. The primary method of learning is backpropagation, in which the network adjusts its internal weights to minimize error and produce meaningful outputs from input data.


2. Neural Network Zoo: Learning with Animal Cards

To make these concepts easier and more fun to learn, the Neural Network Zoo represents each type of neural network as an animal. Each animal metaphor reflects the network's characteristics and strengths:

- Convolutional Neural Network (CNN) → Cheetah: Fast and sharp at spotting visual patterns.
- Long Short-Term Memory (LSTM) → Lemur: Excellent at remembering important sequences.
- Recurrent Neural Network (RNN) → Raccoon: Works well with sequences by feeding information back to itself.

Neural network architectures vary depending on how inputs are propagated. CNNs use convolutional filters to detect spatial patterns in grids (like images). RNNs and LSTMs have feedback loops for sequential memory over time. Transformers utilize self-attention to process sequences simultaneously. Each network type is optimized for specific tasks: CNNs for vision, LSTMs for time-series forecasting, and Transformers for natural language processing.

Structure of a Neural Network




Animal Name: Cheetah

Neural Network Type: Convolutional Neural Network (CNN)

Structure of the Network

CNNs are composed of several specialized layers: convolutional layers, pooling layers, and fully connected layers. Convolutional layers detect features like edges, textures, and shapes. Pooling layers reduce data size while preserving important information. Fully connected layers use extracted features to make predictions or classifications.



How It Works


CNNs scan input images using small filters to detect features. These filters capture simple patterns (edges) in early layers, progressing to complex objects in deeper layers.

Why This Animal Fits

The cheetah represents CNNs due to its speed and precision in detecting patterns, mirroring how CNNs analyze images efficiently.

Typical Applications

CNNs are widely used in image recognition, facial recognition, medical imaging, autonomous vehicles, and security surveillance.




CNNs are computationally efficient due to parameter sharing and sparse connectivity, where the same filters scan the entire input while focusing only on local receptive fields. Convolutional layer behavior depends on hyperparameters such as kernel size, stride, and padding, which influence feature extraction and output dimensions. Pooling layers perform downsampling to reduce computation and provide partial translation invariance, commonly using max pooling or average pooling. Modern CNNs also apply regularization techniques like Dropout and Batch Normalization to improve training stability and reduce overfitting, enabling better generalization in visual recognition tasks.

4. Neural Network Animal Profile: RNN – The Raccoon

Animal Name: Raccoon

Neural Network Type: Recurrent Neural Network (RNN)

RNNs have a feedback loop that allows them to remember information from previous steps, making them ideal for sequential data such as text, speech, or time-series information. This ability to “remember” helps them in tasks like language modeling, speech recognition, and predictive text. The raccoon metaphor reflects the cleverness and memory aspect of RNNs.



In practice, RNNs maintain a hidden state that is updated at each time step, allowing information from previous inputs to influence current predictions. However, standard RNNs often struggle with vanishing and exploding gradient problems, which limit their ability to capture long-term dependencies. To address this limitation, advanced variants such as Long Short-Term Memory (LSTM) networks introduce a structured cell state regulated by input, forget, and output gates, enabling more stable information retention over extended sequences. These mechanisms significantly improve performance in tasks involving long-term temporal relationships, such as language modeling and speech recognition.

5. Zoo Tour Reflection and Comparison

CNNs, RNNs, and LSTMs serve different purposes:

- CNNs → spatial data (images)
- RNNs/LSTMs → sequential data (text, speech, time-series)
- Transformers → attention-based sequence modeling

This illustrates that no single network is universal; networks are selected based on task requirements.

CNN

The Cheetah

Ultra-fast vision! Detects patterns and shapes in the blink of an eye.

Neural Network Zoo

RNN

The Raccoon

Sequence seeker! Remembers what happened just a moment ago.

Neural Network Zoo

LSTM

The Lemur

Long-term legacy! Keeps track of history for ages without forgetting.

Neural Network Zoo

GAN

The Gecko

Master of disguise! Creates realistic fakes to challenge the status quo.

Neural Network Zoo

Transformer

The Toucan

Global attention! Understands context by looking at everything at once.

Neural Network Zoo

Autoencoder

The Armadillo

Squeeze & Shield! Compresses data into a tiny core then reconstructs it.

Neural Network Zoo

MLP

The Meerkat

The Sentinel! A classic group that processes data layer by layer.

Neural Network Zoo

SOM

The Spider

Web Weaver! Organizes chaotic data into a neat 2D map of reality.

Neural Network Zoo

6. Reflection and Deeper Understanding

The Neural Network Zoo metaphor proved to be an invaluable learning tool, as it allowed abstract concepts in AI and deep learning to be visualized in a concrete, memorable way. By associating neural networks with animals, it became easier to remember each network's strengths, limitations, and the type of tasks they excel at. For example, representing CNNs as cheetahs highlighted their speed and precision in detecting spatial patterns, while representing RNNs as raccoons emphasized their ability to remember and process sequences over time.

This exercise also enhanced my understanding of how architectural design directly impacts functionality. I gained insight into how CNNs, RNNs, and LSTMs differ fundamentally in terms of layer structure, data propagation, and memory handling. Visualizing these networks as animals clarified why some networks are suited for images, others for text or time-series data, and why Transformers have become critical for processing sequential information with attention mechanisms.

Moreover, this activity reinforced the importance of task-specific model selection in real-world AI applications. For instance, in autonomous driving, CNNs are deployed for object and obstacle recognition, while RNNs are used in predictive text and speech applications. Understanding these distinctions helps in designing AI solutions that are both effective and efficient. Overall, the zoo metaphor transformed a technical topic into an engaging, intuitive, and practical learning experience.

In addition to understanding architectural differences, it is also important to recognize the limitations and challenges associated with neural networks. One major issue is data dependency: deep learning models typically require large, well-labeled datasets to achieve high performance. Insufficient or imbalanced data can lead to biased model behavior and poor generalization. Furthermore, neural networks are prone to overfitting, where a model

learns training data too closely and fails to perform well on unseen data. Techniques such as regularization, dropout, and batch normalization are commonly used to improve generalization and training stability.

From a practical perspective, neural networks also involve significant computational and resource requirements. Training large-scale models demands substantial processing power, often relying on GPU or distributed computing systems, which increases cost and energy consumption. This highlights the growing importance of efficient model design and sustainable AI practices.

7. Conclusion

The Neural Network Zoo serves as a creative yet effective framework for studying and understanding neural networks. By using animal metaphors, complex network structures and operations become more approachable, which enhances comprehension and retention of deep learning concepts. The CNN Cheetah demonstrates the network's capability for rapid and accurate image recognition, while the RNN Raccoon emphasizes sequential memory and temporal data processing.

This exercise highlighted that neural networks are not one-size-fits-all tools; their performance is highly dependent on architecture design and the type of data they process. Additionally, understanding the strengths and limitations of each network type is essential when applying AI to real-world problems such as facial recognition, speech synthesis, time-series forecasting, and natural language processing.

Finally, the activity also illustrated the interconnectedness of AI concepts. Layers, neurons, memory, and attention mechanisms collaborate to enable intelligent behavior. Learning these concepts through visualization, metaphor, and reflection strengthens the ability to design, implement, and evaluate AI systems responsibly and effectively. This deeper understanding lays a solid foundation for further exploration and innovation in the field of artificial intelligence.